

Reservoir computing

Karol Bednarz

September 14, 2025

Contents

1	Introduction	1
1.1	Principles of reservoir computing	1
1.2	Echo state property	2
2	Learning algorithm	2
3	Sample implementation in Python	4

1 Introduction

1.1 Principles of reservoir computing

The state of reservoir dynamics can be expressed as:

$$h_t = f((1 - k) \cdot u_t \cdot W_{\text{in}} + k \cdot h_{t-1} \cdot W_{\text{h}} + y_{t-1} \cdot W_{\text{1b}} + b) \quad (1)$$

Where:

- h_{t-1} – are the reservoir state respectively, from the previous time step,
- u_t – is the observed data at time step t ,
- y_{t-1} – is the the predicted output state $t - 1$,
- $W_{\text{in}} \in \mathbb{R}^{N_u \times N_h}$ – is the input weight matrix,
- $W_{\text{h}} \in \mathbb{R}^{N_h \times N_h}$ – is the internal weight matrix,
- $W_{\text{1b}} \in \mathbb{R}^{N_h \times N_y}$ – is the output feedback weight matrix,
- $b \in \mathbb{R}^{N_h}$ – is the bias vector.
- f – is the activation function, typically tanh or sigmoid,
- k – is the leaking rate, typically $k \in [0.1, 0.3]$.

With the computed reservoir dynamics, the output can be then obtained by:

$$y_t = h_t \cdot W_{\text{out}} \quad (2)$$

Where:

$W_{\text{out}} \in \mathbb{R}^{N_h \times N_y}$ – is the output weight matrix.

1.2 Echo state property

Any system that changes in a nonlinear way can work as the reservoir. However, starting a nonlinear system with random connection strengths creates problems. The reservoir is a system that feeds its outputs back into itself. This can make it unstable if the connection strengths aren't set up correctly. For example, if the internal connections are too strong, the system might get stuck giving the same output regardless of what input it receives. The random connection strengths must be chosen so the system doesn't grow out of control. For the system to work well, it must follow something called the "echo state property." This rule ensures that the reservoir's behavior eventually depends on the input signal rather than just its starting conditions. To meet this requirement, the internal connection matrix W_h is first set up using random values between -1 and 1. This matrix is then adjusted one time according to the echo state property rule:

$$W'_h = \alpha \odot W_h \quad (3)$$

$$W_h^\dagger = \frac{\rho W_h}{|\lambda_{\max}(W_h)|} \quad (4)$$

Where:

- $\rho \in (0, 1)$ – is the spectral radius, typically $\rho \in [0.9, 1]$
- $\lambda_{\max}(W_h)$ – is the largest eigenvalue of W_h .
- $\alpha \in (0, 1)$ – is the sparsity coefficient, typically $\alpha \in [0.1, 0.3]$.

The spectral radius is a parameter that determines the amount of nonlinear interaction of input components through time.

Due to the recursive nature of the reservoir layer, such dynamics reflect trajectories of the past historical input the short-term memory (known as the fading memory). As another critical property for computing the RC principle, short-term memory can be quantitative measured by the coefficient of memory capacity

$$MC = \sum_{k=1}^{\infty} MC_k = \sum_{k=1}^{\infty} d^2(u_{t-k}, y_t) = \sum_{k=1}^{\infty} \frac{\text{cov}^2(u_{t-k}, y_t)}{\sigma^2(u_t)\sigma^2(y_t)} \quad (5)$$

Where:

- $d^2(u_{t-k}, y_t)$ – is the square of the correlation coefficient between the output y_t and the input u_{t-k} with a delay of k time steps,

According to the Lyapunov stability analysis, a large memory capacity is needed to compute the RC principle, which can be achieved at the asymptotically stable region.

2 Learning algorithm

The training of the reservoir computing model involves adjusting only the output weights W_{out} . The input weights W_{in} , internal weights W_h , and feedback weights W_{1b} are typically initialized randomly and remain fixed during training. The training process can be summarized in the following steps:

$$Y = H \cdot W_{\text{out}} \quad (6)$$

Where:

- $Y \in \mathbb{R}^{T \times N_y}$ – is the matrix of target outputs for all time steps,

$H \in \mathbb{R}^{T \times N_h}$ – is the matrix of reservoir states for all time steps,
 T – is the total number of time steps.

In general, the W_{out} can be directly obtained by calculating the Moore-Penrose pseudoinverse of the reservoir states matrix H with respect to the target outputs matrix Y :

$$W_{\text{out}} = Y \cdot H^\dagger \cdot (H \cdot H^\dagger + \eta I)^{-1} \quad (7)$$

Where:

H^\dagger – is the Moore-Penrose pseudoinverse of matrix H ,
 η – is the regularization parameter, typically $\eta \in [10^{-6}, 10^{-2}]$,
 I – is the identity matrix of size $N_h \times N_h$.

3 Sample implementation in Python

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import torch
4 import torch.nn as nn
5 from scipy.integrate import solve_ivp
6
7
8 class LorenzESN(nn.Module):
9     """
10     Echo State Network with Lorenz system dynamics in the reservoir.
11
12     The reservoir states evolve according to modified Lorenz equations,
13     providing rich chaotic dynamics for temporal pattern learning.
14     """
15
16     def __init__(
17         self,
18         input_size,
19         reservoir_size,
20         output_size,
21         spectral_radius=0.9,
22         input_scaling=1.0,
23         leaking_rate=1.0,
24         lorenz_coupling=0.1,
25         sigma=10.0,
26         rho=28.0,
27         beta=8.0 / 3.0,
28     ):
29         """
30         Args:
31             input_size: Dimension of input
32             reservoir_size: Size of reservoir (should be multiple of 3 for Lorenz
33 components)
34             output_size: Dimension of output
35             spectral_radius: Spectral radius of reservoir weight matrix
36             input_scaling: Scaling factor for input weights
37             leaking_rate: Leaking rate for reservoir updates
38             lorenz_coupling: Coupling strength between Lorenz subsystems
39             sigma, rho, beta: Lorenz system parameters
40         """
41         super(LorenzESN, self).__init__()
42
43         self.input_size = input_size
44         self.reservoir_size = reservoir_size
45         self.output_size = output_size
46         self.leaking_rate = leaking_rate
47         self.lorenz_coupling = lorenz_coupling
48
49         # Lorenz parameters
50         self.sigma = sigma
51         self.rho = rho
52         self.beta = beta
53
54         # Ensure reservoir size is multiple of 3 for Lorenz triplets
55         self.num_lorenz_systems = reservoir_size // 3
56         self.actual_reservoir_size = self.num_lorenz_systems * 3
57
58         # Initialize input weights
59         self.W_in = nn.Parameter(
```

```

59         torch.randn(self.actual_reservoir_size, input_size) * input_scaling,
60         requires_grad=False,
61     )
62
63     # Initialize reservoir coupling weights (sparse connectivity between Lorenz
64     systems)
65     W_res = torch.randn(self.actual_reservoir_size, self.actual_reservoir_size)
66     W_res = self.make_sparse(W_res, sparsity=0.1) # 10% connectivity
67
68     # Scale to desired spectral radius
69     eigenvals = torch.linalg.eigvals(W_res)
70     current_spectral_radius = torch.max(torch.abs(eigenvals)).item()
71     W_res = W_res * (spectral_radius / current_spectral_radius)
72
73     self.W_res = nn.Parameter(W_res, requires_grad=False)
74
75     # Output weights (trainable)
76     self.W_out = nn.Linear(self.actual_reservoir_size, output_size)
77
78     # Initialize reservoir state
79     self.register_buffer("reservoir_state", torch.zeros(self.actual_reservoir_size))
80
81     def make_sparse(self, matrix, sparsity=0.1):
82         """Make matrix sparse by randomly setting elements to zero"""
83         mask = torch.rand_like(matrix) < sparsity
84         return matrix * mask.float()
85
86     def lorenz_derivatives(self, state):
87         """Compute Lorenz derivatives for the entire reservoir state"""
88         # Reshape to (num_systems, 3) for easier processing
89         lorenz_state = state.view(self.num_lorenz_systems, 3)
90
91         # Compute Lorenz derivatives for each system
92         x, y, z = lorenz_state[:, 0], lorenz_state[:, 1], lorenz_state[:, 2]
93
94         dx_dt = self.sigma * (y - x)
95         dy_dt = x * (self.rho - z) - y
96         dz_dt = x * y - self.beta * z
97
98         # Stack derivatives
99         derivatives = torch.stack([dx_dt, dy_dt, dz_dt], dim=1)
100
101         return derivatives.view(-1) # Flatten back to 1D
102
103     def lorenz_dynamics(self, state, dt=0.001):
104         """
105         Apply Lorenz dynamics to reservoir state.
106         State is organized as [x1,y1,z1, x2,y2,z2, ..., xN,yN,zN]
107         """
108         # TODO: Apply RK4 integration instead of Euler
109         k1 = dt * self.lorenz_derivatives(state)
110         k2 = dt * self.lorenz_derivatives(state + 0.5 * k1)
111         k3 = dt * self.lorenz_derivatives(state + 0.5 * k2)
112         k4 = dt * self.lorenz_derivatives(state + k3)
113         new_state = state + (k1 + 2 * k2 + 2 * k3 + k4) / 6
114         return new_state
115
116     def forward(self, input_sequence):
117         """
118         Forward pass through the ESN.
119
120         Args:

```

```

120         input_sequence: (seq_len, batch_size, input_size) or (seq_len, input_size)
121
122     Returns:
123         outputs: (seq_len, batch_size, output_size) or (seq_len, output_size)
124     """
125     if input_sequence.dim() == 2:
126         input_sequence = input_sequence.unsqueeze(1) # Add batch dimension
127         squeeze_output = True
128     else:
129         squeeze_output = False
130
131     seq_len, batch_size, _ = input_sequence.shape
132
133     # Initialize states for batch
134     reservoir_states = self.reservoir_state.unsqueeze(0).repeat(batch_size, 1)
135     all_states = []
136
137     for t in range(seq_len):
138         # Current input
139         current_input = input_sequence[t] # (batch_size, input_size)
140
141         for b in range(batch_size):
142             # Apply Lorenz dynamics
143             reservoir_states[b] = self.lorenz_dynamics(reservoir_states[b])
144
145             # Add input and reservoir coupling
146             input_contribution = torch.matmul(self.W_in, current_input[b])
147             reservoir_contribution = torch.matmul(self.W_res, reservoir_states[b])
148
149             # Leaky integration
150             new_state = (1 - self.leaking_rate) * reservoir_states[
151                 b
152             ] + self.leaking_rate * torch.tanh(
153                 input_contribution + self.lorenz_coupling * reservoir_contribution
154             )
155
156             reservoir_states[b] = new_state
157
158         all_states.append(reservoir_states.clone())
159
160     # Stack all states: (seq_len, batch_size, reservoir_size)
161     all_states = torch.stack(all_states)
162
163     # Compute outputs
164     outputs = self.W_out(all_states)
165
166     if squeeze_output:
167         outputs = outputs.squeeze(1)
168
169     return outputs
170
171     def reset_state(self):
172         """Reset reservoir state"""
173         self.reservoir_state.zero_()
174
175
176 def generate_lorenz_data(num_steps=1000, dt=0.01, sigma=10.0, rho=28, beta=8.0 / 3.0):
177     """Generate Lorenz attractor data for testing"""
178
179     def lorenz(t, state):
180         x, y, z = state
181         return [sigma * (y - x), x * (rho - z) - y, x * y - beta * z]

```

```

182
183 t_span = (0, num_steps * dt)
184 t_eval = np.arange(0, num_steps * dt, dt)
185 initial_state = [0.1, 0.1, 0.1]
186
187 sol = solve_ivp(lorenz, t_span, initial_state, t_eval=t_eval, method="DOP853")
188 return torch.tensor(sol.y.T, dtype=torch.float32) # Shape: (num_steps, 3)
189
190
191 def train_esn_example():
192     """Example training script"""
193
194     # Generate synthetic data (predicting next step of Lorenz system)
195     data = generate_lorenz_data(num_steps=2000, dt=0.01)
196     # data =
197
198     # Prepare sequences
199     seq_length = 100
200     X = []
201     y = []
202
203     for i in range(len(data) - seq_length):
204         X.append(data[i : i + seq_length])
205         y.append(data[i + 1 : i + seq_length + 1])
206
207     X = torch.stack(X) # (num_sequences, seq_length, 3)
208     y = torch.stack(y) # (num_sequences, seq_length, 3)
209
210     # Split data
211     train_size = int(0.8 * len(X))
212     X_train, X_test = X[:train_size], X[train_size:]
213     y_train, y_test = y[:train_size], y[train_size:]
214
215     # Create ESN
216     esn = LorenzESN(
217         input_size=3,
218         reservoir_size=500 * 3, # 100 Lorenz systems
219         output_size=3,
220         spectral_radius=0.95,
221         input_scaling=1.0,
222         leaking_rate=0.1,
223         lorenz_coupling=0.05,
224     )
225
226     print("Training ESN with Lorenz reservoir dynamics...")
227
228     # Training (only train output weights using ridge regression)
229     esn.eval() # Set to eval mode to disable gradient computation for reservoir
230
231     # Collect reservoir states for all training sequences
232     all_reservoir_states = []
233     all_targets = []
234
235     with torch.no_grad():
236         for i in range(len(X_train)):
237             esn.reset_state()
238             states = esn.forward(X_train[i]) # Don't use outputs, just collect states
239
240             # Get intermediate states from the reservoir
241             reservoir_states = []
242             esn.reset_state()
243             for t in range(X_train[i].shape[0]):

```

```

244         # Manual forward pass to collect states
245         current_input = X_train[i][t]
246         esn.reservoir_state = esn.lorenz_dynamics(esn.reservoir_state)
247
248         input_contribution = torch.matmul(esn.W_in, current_input)
249         reservoir_contribution = torch.matmul(esn.W_res, esn.reservoir_state)
250
251         new_state = (
252             1 - esn.leaking_rate
253         ) * esn.reservoir_state + esn.leaking_rate * torch.tanh(
254             input_contribution + esn.lorenz_coupling * reservoir_contribution
255         )
256
257         esn.reservoir_state = new_state
258         reservoir_states.append(esn.reservoir_state.clone())
259
260         reservoir_states = torch.stack(reservoir_states)
261         all_reservoir_states.append(reservoir_states)
262         all_targets.append(y_train[i])
263
264     # Concatenate all data
265     X_reservoir = torch.cat(
266         all_reservoir_states, dim=0
267     ) # (total_timesteps, reservoir_size)
268     y_flat = torch.cat(all_targets, dim=0) # (total_timesteps, 3)
269
270     # Ridge regression for output weights
271     ridge_param = 1e-6
272     I = torch.eye(X_reservoir.shape[1])
273
274     # Solve:  $W_{out} = (X^T X + I)^{-1} X^T y$ 
275     XTX = torch.matmul(X_reservoir.T, X_reservoir)
276     XTy = torch.matmul(X_reservoir.T, y_flat)
277     W_out_optimal = torch.linalg.solve(XTX + ridge_param * I, XTy)
278
279     # Set the optimal weights
280     esn.W_out.weight.data = W_out_optimal.T
281     esn.W_out.bias.data.zero_()
282
283     # Test the model
284     test_predictions = []
285     test_targets = []
286
287     with torch.no_grad():
288         for i in range(min(5, len(X_test))): # Test on first 5 sequences
289             esn.reset_state()
290             pred = esn.forward(X_test[i])
291             test_predictions.append(pred.numpy())
292             test_targets.append(y_test[i].numpy())
293
294     # Calculate MSE
295     mse = np.mean(
296         [(pred - target) ** 2 for pred, target in zip(test_predictions, test_targets)]
297     )
298     print(f"Test MSE: {mse:.6f}")
299
300     return esn, test_predictions, test_targets, data.numpy()
301
302
303 if __name__ == "__main__":
304     # Run example
305     esn, predictions, targets, original_data = train_esn_example()

```



```

306
307 # Plot results
308 plt.figure(figsize=(15, 10))
309
310 # Plot original Lorenz attractor
311 plt.subplot(2, 3, 1)
312 plt.plot(original_data[:1000, 0], original_data[:1000, 2])
313 plt.title("Original Lorenz Attractor (X-Z plane)")
314 plt.xlabel("X")
315 plt.ylabel("Z")
316
317 # Plot prediction vs target for first test sequence
318 if predictions:
319     pred = predictions[0]
320     target = targets[0]
321
322     plt.subplot(2, 3, 2)
323     plt.plot(target[:, 0], "b-", label="Target X", alpha=0.7)
324     plt.plot(pred[:, 0], "r--", label="Predicted X", alpha=0.7)
325     plt.title("X Component Prediction")
326     plt.legend()
327
328     plt.subplot(2, 3, 3)
329     plt.plot(target[:, 1], "b-", label="Target Y", alpha=0.7)
330     plt.plot(pred[:, 1], "r--", label="Predicted Y", alpha=0.7)
331     plt.title("Y Component Prediction")
332     plt.legend()
333
334     plt.subplot(2, 3, 4)
335     plt.plot(target[:, 2], "b-", label="Target Z", alpha=0.7)
336     plt.plot(pred[:, 2], "r--", label="Predicted Z", alpha=0.7)
337     plt.title("Z Component Prediction")
338     plt.legend()
339
340 # 3D phase space comparison
341 plt.subplot(2, 3, 5)
342 plt.plot(target[:, 0], target[:, 2], "b-", label="Target", alpha=0.7)
343 plt.plot(pred[:, 0], pred[:, 2], "r--", label="Predicted", alpha=0.7)
344 plt.title("Phase Space (X-Z)")
345 plt.xlabel("X")
346 plt.ylabel("Z")
347 plt.legend()
348
349 # Error plot
350 plt.subplot(2, 3, 6)
351 error = np.abs(pred - target)
352 plt.plot(error[:, 0], label="X error")
353 plt.plot(error[:, 1], label="Y error")
354 plt.plot(error[:, 2], label="Z error")
355 plt.title("Absolute Error")
356 plt.legend()
357
358 plt.tight_layout()
359 plt.show()

```